

## Experiences with Effects in OCaml 5.0

Anil Madhavapeddy (speaker)   Thomas Leonard(speaker)  
Craig Ferguson   Patrick Ferris   Sadiq Jaffer   Tom Kelly  
KC Sivaramakrishnan

University of Cambridge and Tarides

WASM Stack Switching WG, Feb 2022

# Overview

- ▶ Background: OCaml 5.0
- ▶ Introduction to effects
- ▶ Case study: Converting the Angstrom parser
- ▶ Eio concurrency library

## Background: OCaml and the road to 5.0

- ▶ Industrial-grade functional programming language, first released in 1996 and continuously developed since then.
- ▶ Compiles native code binaries for x86/arm/ppc/riscv in 32- and 64-bit.
- ▶ Also has a portable bytecode compiler that just needs a C compiler, and can be compiled to JavaScript (`js_of_ocaml`).
- ▶ OCaml 5.0 will feature multicore parallelism, and also untyped effects.
  - ▶ **Pros:** High-performance direct-style code with a GC (this talk)
  - ▶ **Cons:** How do retain portability to JavaScript and Wasm?

## Background: asynchrony in OCaml 4 and earlier

- ▶ OCaml 4 is single-threaded with no first-class support for concurrency.
- ▶ IO concurrency has been expressed for years via userlevel libraries that allow for futures to be expressed succinctly.
- ▶ Two widely adopted libraries are:
  - ▶ Lwt (usually for web programming)
  - ▶ Async (used by Jane Street in their production usage of OCaml)

## OCaml 4.0 Lwt example

```
let foo ~stdin total =  
  Lwt_io.read_line stdin >>= fun →  
  Lwt_io.printf "n/total = %d"  
    (int_of_string n / total)
```

Fatal error: exception Division\_by\_zero  
Raised at Lwt\_example.foo in file "lwt\_example.ml", line 6  
Called from Lwt.[...].callback in file "src/core/lwt.ml", ...

- ▶ Backtrace doesn't say what called foo
- ▶ Closure with total allocated on the heap
- ▶ Type of function foo appends an Lwt.t

## OCaml 5.0 effects-based example

```
let foo ~stdin total =  
  let n = read_line stdin in  
  traceLn "n/total = %d"  
    (int_of_string n / total)
```

Fatal error: exception Division\_by\_zero

Raised at Eio\_example.foo in file "eio\_example.ml", line 11

Called from Eio\_example.bar in file "eio\_example.ml", line 15

...

- ▶ Backtrace is entirely accurate now
- ▶ Only stack allocation needed for the blocking I/O
- ▶ Type of function is no longer affected by use of IO

# Introduction to effects

- ▶ Resumable exceptions
- ▶ Multiple stacks

```
effect Foo : int → int
```

```
try
```

```
  println "step 1";
```

```
  let x = perform (Foo 2) in
```

```
  println "step %d" x
```

```
with effect (Foo n) k →
```

```
  println "step %d" n;
```

```
  continue k (n + 1)
```

# Advantages of effects

- ▶ No difference between sequential and concurrent code.
  - ▶ No special monad syntax.
  - ▶ Can use `try`, `match`, `while`, etc.
  - ▶ No separate `lwt` or `async` versions of code.
- ▶ No heap allocations needed to simulate a stack.
- ▶ A real stack means backtraces and profiling tools work.



## Case study: Angstrom

<https://github.com/inhabitedtype/angstrom/>

- ▶ A library for writing parsers
- ▶ Designed for network protocols
- ▶ Strong focus on performance

## A toy parser

```
type 'a parser = state → 'a
```

```
let any_char state =  
  ensure 1 state;  
  let c = Input.unsafe_get_char state.input state.pos in  
  state.pos <- state.pos + 1;  
  c
```

```
let (*>) a b state =  
  let _ = a state in  
  b state
```

# The Angstrom parser type

```
module State = struct
  type 'a t =
    | Partial of 'a partial
    | Lazy     of 'a t Lazy.t
    | Done     of int * 'a
    | Fail     of int * string list * string
  and 'a partial =
    { committed : int;
      continue  : Bigstringaf.t →
        off:int → len:int → More.t → 'a t }
end
type 'a with_state = Input.t → int → More.t → 'a
type 'a failure =
  (string list → string → 'a State.t) with_state
type ('a, 'r) success = ('a → 'r State.t) with_state
type 'a parser = { run : 'r.
  ('r failure → ('a, 'r) success → 'r State.t) with_state
}
```

# Angstrom parsers

```
let any_char =  
  ensure 1 { run = fun input pos more _fail succ →  
    succ input (pos + 1) more  
    (Input.unsafe_get_char input pos)  
  }
```

```
let (*>) a b =  
  { run = fun input pos more fail succ →  
    let succ' input' pos' more' _ =  
      b.run input' pos' more' fail succ in  
    a.run input pos more fail succ'  
  }
```

## Angstrom : effects branch

<https://github.com/talex5/angstrom/tree/effects>

```
type 'a parser = state → 'a
```

```
let any_char state =  
  ensure 1 state;  
  let c = Input.unsafe_get_char state.input state.pos in  
  state.pos <- state.pos + 1;  
  c
```

```
let (*>) a b state =  
  let _ = a state in  
  b state
```

## Parser micro-benchmark

```
let parser = skip_many any_char
```

	Time	MinWrds	MajWrds
Callbacks	750.63ms	160.04Mw	8,9944.00kw
Effects	57.81ms	-	-

13 times faster!

## Parser micro-benchmark

```
let parser = skip_many any_char
```

	Time	MinWrds	MajWrds
Callbacks	750.63ms	160.04Mw	8,9944.00kw
Callbacks'	180.73ms	220.01Mw	9,659.00w
Effects	57.81ms	-	-

3 times faster!

## Realistic parser benchmark

Parsing an HTTP request shows smaller gains:

	Time	MinWrds	MajWrds
Callbacks	60.30ms	9.28Mw	102.08kw
Effects	50.71ms	2.13Mw	606.30w



## Using effects for backwards compatibility

```
effect Read : int → state
let read c = perform (Read c)

let parse p =
  let buffering = Buffering.create () in
  try Unbuffered.parse ~read p
  with effect (Read committed) k →
    Buffering.shift buffering committed;
    Partial (fun input →
      Buffering.feed_input buffering input;
      continue k (Buffering.for_reading buffering)
    )
```

(simplified)

## Angstrom summary

- ▶ Slightly faster
- ▶ Much simpler code
- ▶ No effects in interface
- ▶ Can convert between callbacks and effects easily


## Eio : an IO library using effects for concurrency

- ▶ Alternative to Lwt and Async
- ▶ Generic API that performs effects
- ▶ Cross-platform libuv effect handler
- ▶ High-performance io-uring handler for Linux

## Eio example

```
let handle_connection =  
  Httpaf_eio.Server.create_connection_handler  
    ~config  
    ~request_handler  
    ~error_handler  
  
let main ~net =  
  Switch.top @@ fun sw →  
    let socket = Eio.Net.listen ~sw net ('Tcp (host, port))  
      ~reuse_addr:true  
      ~backlog:1000  
    in  
    while true do  
      Eio.Net.accept_sub ~sw socket handle_connection  
        ~on_error:log_connection_error  
    done
```

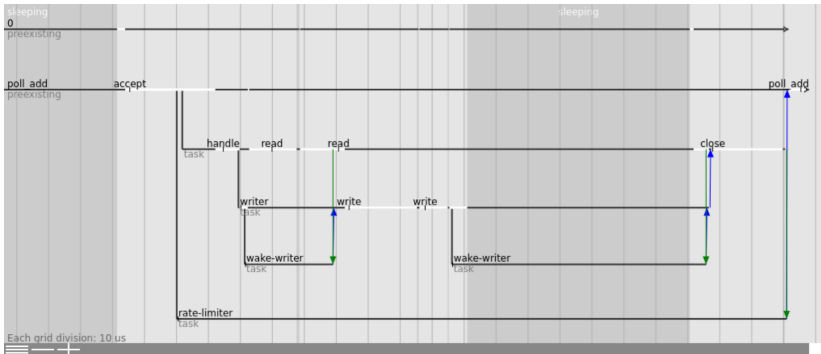
## HTTP benchmark



100 concurrent connections. Servers limited to 1 core.

# Eio : other features

- ▶ Structured concurrency
- ▶ OCaps security model
- ▶ Tracing support
- ▶ Supports multiple cores
- ▶ Still experimental

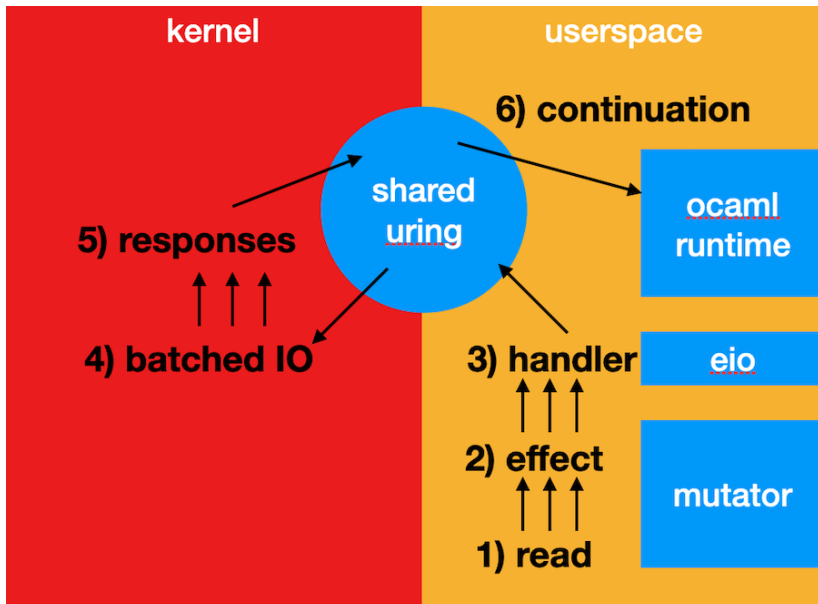


## Eio : migrating from old-style code

Switching to effects in OCaml 5.0 turns out to be great timing in the bigger picture.

- ▶ There has been a slow but steady shift to better system interfaces for async
- ▶ `io_uring` (Linux), Grand Central Dispatch (macOS), IOCP (Windows)
- ▶ These all make effect-based IO incredibly straightforward and elegant.

## Eio : post-POSIX, with io\_uring





## Eio : post-POSIX

- ▶ No more fd-set management and scalability bottlenecks, so great time for new post-POSIX interfaces
- ▶ **Concurrency-friendly:** Just stash a single-shot continuation and call it when IO is ready, or raise exception if IO is cancelled.
- ▶ **Parallel-friendly:** Push batch onto a shared memory ring and get responses back with one syscall.
- ▶ **Hardware-friendly:** Very similar to hypervisor-level interfaces, but from userspace.

# Summary of OCaml 5.0 and our use of effects

- ▶ Concurrency with effects works very well and is ergonomic to program with
- ▶ Effects have very good performance (stack vs heap)
- ▶ The use of separate of effect schedulers is still emerging, but there are dozens of networking/storage OCaml libraries being ported currently, with little drama.
- ▶ Key open blocker for our community is Js/Wasm compilation support: **effects are here to stay in OCaml 5.0, so what's the best path forward?**

<https://github.com/ocaml-multicore/eio> documentation shows how to try out OCaml effects.

<https://github.com/patricoferris/awesome-multicore-ocaml> lists community libraries.